

TDS3651

Visual Information Processing



Filtering Lecture 3

Faculty of Computing and Informatics

Multimedia University

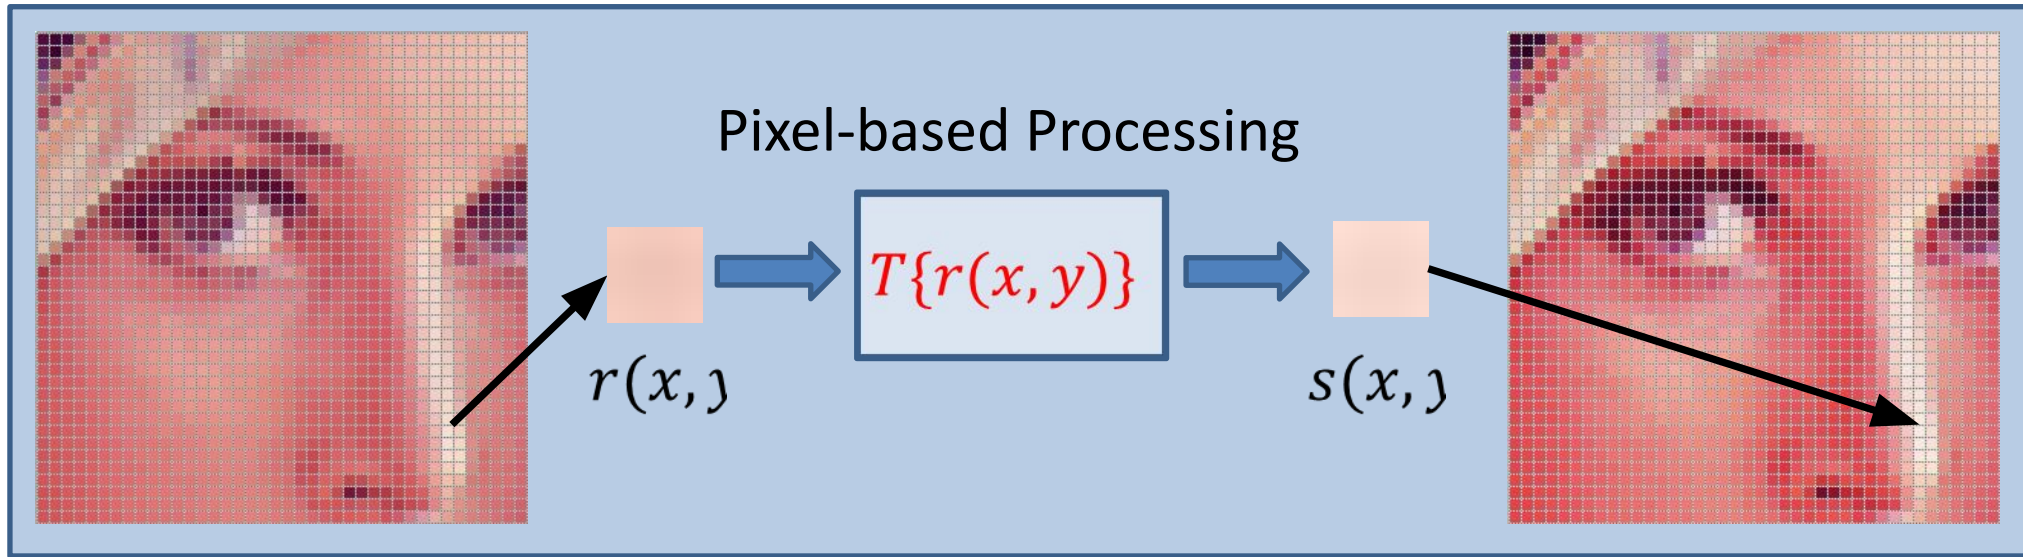
prepared by Lai-Kuan, Wong

modified by Yuen Peng, Loh


Lecture Outline

- Image Filtering
- Correlation and Convolution
- Image Sharpening
- Non-linear filters

Previously



This lecture:
Neighborhood-based
Processing



Applications of Filtering


De-noising



Salt and pepper noise



Super-resolution




Moving beyond pixel \rightarrow pixel

- What if there is insufficient knowledge of how to transform a pixel?


Depending on what we want to achieve, we can use **neighboring pixels** to help provide more information

Motivation: Noise reduction

- Capturing multiple images of the same static scene will not result in identical images
 - Likely: Environmental changes, sensor noise, camera shake, etc.



Motivation: Noise reduction



- Sometimes noises can be introduced during transmission of signals or compression of data

Common types of noise

- **Salt and pepper noise:** random occurrences of black and white pixels
- **Impulse noise:** random occurrences of white pixels
- **Gaussian noise:** variations in intensity drawn from a Gaussian normal distribution



Original



Salt and pepper noise





Impulse noise



Gaussian noise


Neighborhood Processing a.k.a. Filtering



$f(x,y)$

- Idea: Modify the value of the pixel $f(x,y)$ based on a small neighbourhood of pixels surrounding it
- If we wish to “soften” the noise in the image, how should we modify?
 - A. Get minimum pixel
 - B. Get maximum pixel
 - C. Get average pixel
 - D. Use a preset value

To try solve this...



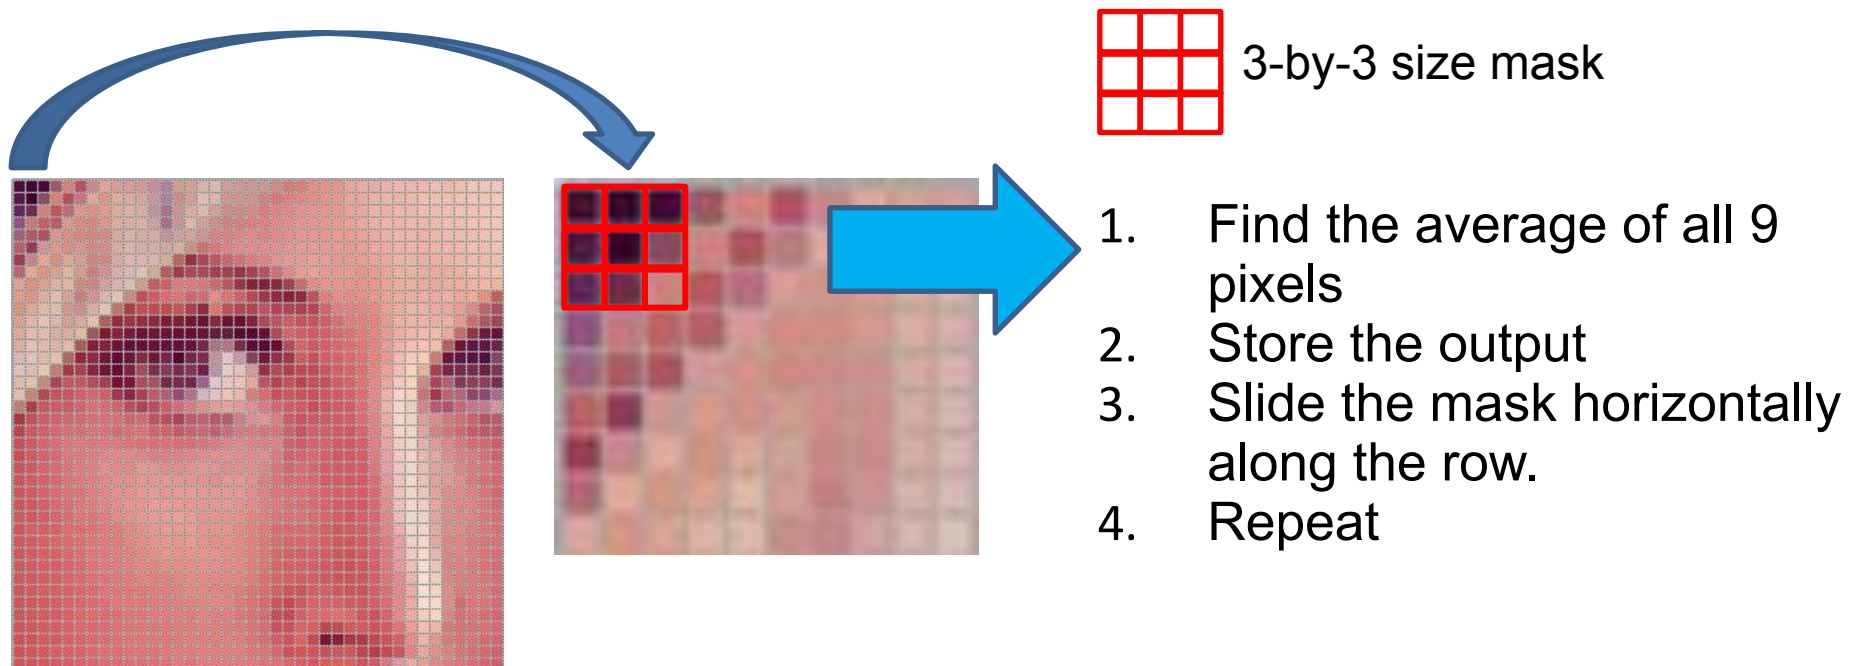
- Idea: Let's replace each pixel with an average of all the values in its neighbourhood
- Assumptions:
 - Expect pixels to be “like” their neighbours
 - Expect noise processes to be independent from pixel to pixel

Result



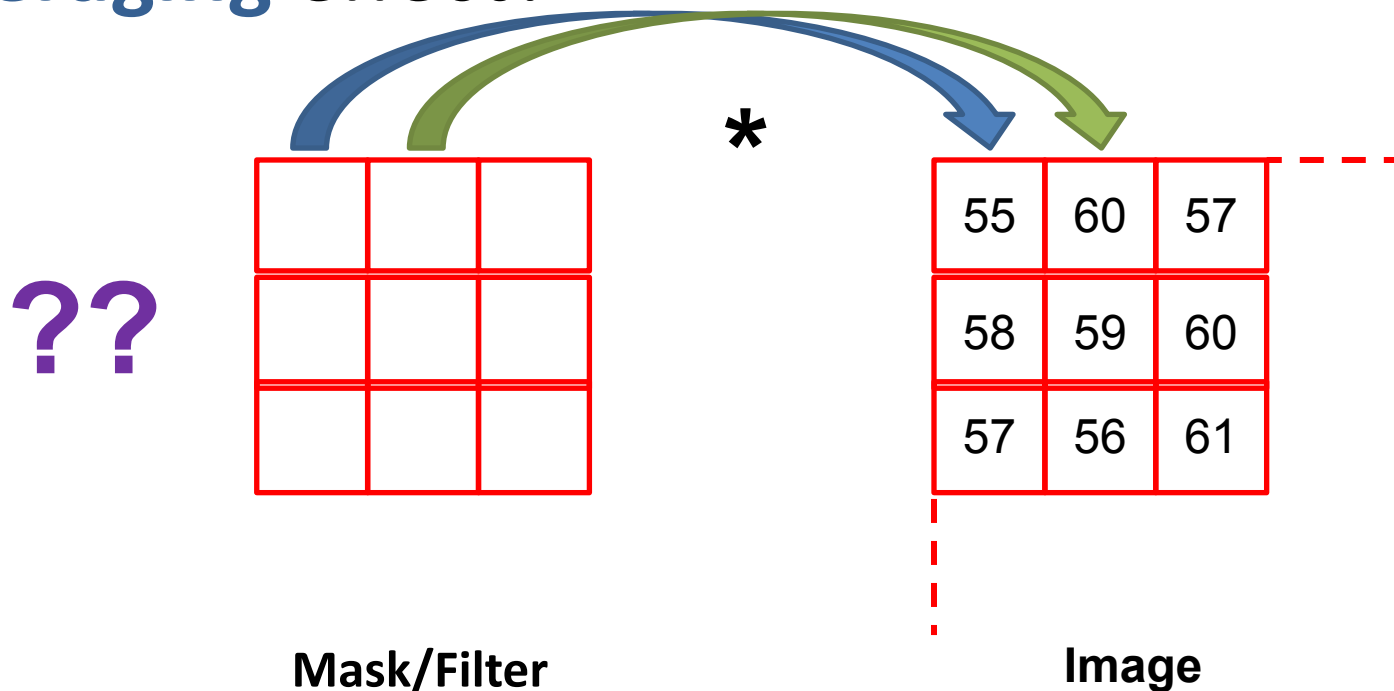
“Masking” or Filtering

- Run a “mask” or “filter” across the entire image
- Mask corresponds to the neighbourhood that we wish to process




“Masking” or Filtering

- Let's say we multiply element-by-element the 3x3 mask/filter against a 3x3 part of the image, what are the values in the mask to achieve the **averaging** effect?




Let's take a look in 1-D

- Let's replace each pixel with an average of all the values in its neighbourhood (also called “moving average”)
- 1D example




Moving average

- Let's replace each pixel with an average of all the values in its **neighbourhood**
- What's this neighbourhood?




Moving average

- Moving average has equal uniform weights in the neighborhood
- Example: $[1\ 1\ 1\ 1\ 1] / 5$



Weighted moving average

- Moving average with non-uniform weights
- Example: $[1 \ 4 \ 6 \ 4 \ 1] / 16$



Moving average in 2D

$$F[x, y]$$
[illegible]
$$G[x, y]$$
[illegible]

Moving average in 2D

$$F[x, y]$$
[illegible]
$$G[x, y]$$
[illegible]

Moving average in 2D

$$F[x, y]$$
[illegible]
$$G[x, y]$$
[illegible]

Moving average in 2D

$$F[x, y]$$

[illegible]

$$G[x, y]$$

[illegible]

Finish: Moving average in 2D

$$F[x, y]$$

[illegible]

$$G[x, y]$$

[illegible]

Finish: Moving average in 2D

$$F[x, y]$$

[illegible]

$$G[x, y]$$

[illegible]

Correlation and Convolution

Correlation filtering

- Say the averaging window size is $2k + 1 \times 2k + 1$:

$$G[i, j] = \underbrace{\frac{1}{(2k + 1)^2}}_{\text{Attribute uniform weight to each pixel}} \underbrace{\sum_{u=-k}^k \sum_{v=-k}^k F[i + u, j + v]}_{\text{Loop over all pixels in neighborhood around image pixel } F[i, j]}$$


- Generalize to allow **different weights** depending on neighbouring pixel's relative position:

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k \underbrace{H[u, v]}_{\text{Non-uniform weights}} F[i + u, j + v]$$

Correlation filtering

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

Move v (column) first



Steps:

1. Filter size = $(2k + 1) \times (2k + 1)$

If filter size = 3×3 , then $k = 1$

2. When $i = 2, j = 9, F[2, 9]$ is the center pixel

From $u = -k = -1, v = -k = -1$


$$F[i + u, j + u] = F[2 - 1, 9 - 1] = F[1, 8]$$

3. Multiply with corresponding filter value

4. Repeat until $u = k = 1, v = k = 1$

$$F[i + u, j + u] = F[2 + 1, 9 + 1] = F[3, 10]$$

5. Sum all values and place in $G[i, j] = G[2, 9]$



$$= (0 \times 1) + (0 \times 2) + (0 \times 3) + (0 \times 4) + (0 \times 5) + (0 \times 6) + (90 \times 7) + (0 \times 8) + (0 \times 9)$$

$$= 630 = G[2, 9]$$


Correlation filtering

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

- Cross-correlation: $G = H \otimes F$
- Summary:
 - Filtering an image: Replace each pixel with a linear combination of its neighbors
 - The filter “kernel” or “mask” $H(u, v)$ is the prescription for the weights in the linear combination

Correlation filtering

- What values belong in the kernel H for the moving average example?



Filter #1: Moving Average




Correlation filtering



depicts box filter:
white = high value, black = low value



original




filtered

- What can we expect from the output if the filter size is 5×5 or 7×7 instead of 3×3 ?

Boundary issues

- What about near the edge?
 - Filter window falls off the edge of the input image, some pixels are not defined \Rightarrow Need to extrapolate
 - Some common padding methods:
 - Constant value (with 0's we get zero-padding)
 - Wrap around
 - Copy edge
 - `numpy.pad` has a lot more options!



Gaussian filter

- What if we want the nearest neighbouring pixels to have **more** influence on the output?

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0


$F[x, y]$

1	2	1
2	4	2
1	2	1


$H[u, v]$

This kernel is an approximation of a 2d Gaussian function:

$$h(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{\sigma^2}}$$




Smoothing with Gaussian filter




Gaussian filter parameters

- What parameters are important?
- **(1) Size of kernel or mask**



$\sigma = 5$ with
10 x 10
kernel




$\sigma = 5$ with
30 x 30
kernel


- Note: Gaussian function has infinite support, but discrete filters use finite kernels

Gaussian filter parameters

- What parameters are important?
- **(2) Variance** of Gaussian: determines extent of smoothing





Size = 30×30
 $\sigma = 2$




Size = 30×30
 $\sigma = 5$

σ parameter

- The σ parameter is the “scale” or “width” or “spread” of the Gaussian kernel \Rightarrow controls the amount of smoothing



...



Properties of Smoothing

- Values are **positive**
- **Sum to 1** \Rightarrow constant regions same as input
- Amount of smoothing proportional to mask size
- “Low-pass” filtering, remove “high-frequency” components

Filtering an impulse signal

- What is the result of filtering the impulse signal (image) F with an arbitrary kernel H ?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$F[x, y]$



a	b	c
d	e	f
g	h	i

$H[u, v]$

$G[x, y]$

Filtering an impulse signal

- What is the result of filtering the impulse signal (image) F with an arbitrary kernel H ?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

$F[x, y]$



a	b	c
d	e	f
g	h	i

$H[u, v]$

= ?

A.

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	a	b	c	0	0
0	0	d	e	f	0	0
0	0	g	h	i	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0

B.

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	i	h	g	0	0
0	0	f	e	d	0	0
0	0	c	b	a	0	0
0	0	0	0	0	0	0
0	0	0	0	0	0	0


Convolution

- Convolution:
 - Flip the filter in both dimensions
(bottom to top, right to left)
 - Then apply cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$

↑
*Notation for
convolution
operator*




Convolution


$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

Steps:


1. Filter size = $(2k + 1) \times (2k + 1)$
If filter size = 3×3 , then $k = 1$
2. When $i = 2, j = 9, F[2,9]$ is the center pixel
From $u = -k = -1, v = -k = -1$
 $F[i - u, j - u] = F[2 + 1, 9 + 1] = F[3,10]$
3. Multiply with corresponding filter value
4. Repeat until $u = k = 1, v = k = 1$
 $F[i - u, j - u] = F[2 + 1, 9 + 1] = F[1,8]$
5. Sum all values and place in $G[i, j] = G[2,9]$




*



To simplify, the **convolution** can be converted to **correlation** by **flipping the filter horizontally and vertically**



\otimes



$$= (0 \times 9) + (0 \times 8) + (0 \times 7) + (0 \times 6) + (0 \times 5) + (0 \times 4) + (90 \times 3) + (0 \times 2) + (0 \times 1) = 270 = G[2,9]$$

Convolution vs. Cross-correlation

Convolution

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$

Cross-correlation

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i + u, j + v]$$

$$G = H \otimes F$$

Important! If the kernel is **symmetric**, **Convolution = Correlation**

Convolution is useful when dealing in the frequency domain (Fourier transform) to enable easy combination of more than 1 filter. Nice explanation on the differences: <http://www.cs.umd.edu/~djacobs/CMSC426/Convolution.pdf>

Separability of Filters

- In some cases, filters are **separable** \Rightarrow can be factored into two steps
 - Convolve all rows
 - Convolve all column

2D convolution (center location only)

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix}$$

The filter factors into a product of 1D filters:

$$\begin{bmatrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{bmatrix} = \begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} \times \begin{bmatrix} 1 & 2 & 1 \end{bmatrix}$$

Perform convolution along rows:

$$\begin{bmatrix} 1 & 2 & 1 \end{bmatrix} * \begin{bmatrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{bmatrix} = \begin{bmatrix} & 11 & & \\ & 18 & & \\ & 18 & & \end{bmatrix}$$

Followed by convolution along the remaining column:

$$\begin{bmatrix} 1 \\ 2 \\ 1 \end{bmatrix} * \begin{bmatrix} & 11 & & \\ & 18 & & \\ & 18 & & \end{bmatrix} = \begin{bmatrix} & & & \\ & & 65 & \\ & & & \end{bmatrix}$$

Separability of the Gaussian filter

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp \left(-\frac{x^2 + y^2}{2\sigma^2} \right) \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp \left(-\frac{x^2}{2\sigma^2} \right) \right) \left(\frac{1}{\sqrt{2\pi}\sigma} \exp \left(-\frac{y^2}{2\sigma^2} \right) \right) \end{aligned}$$

Separability of Filters

- What is the complexity of filtering an $n \times n$ image with an $m \times m$ kernel?

– $O(n^2m^2)$

1	2	1
2	4	2
1	2	1

 \ast

2	3	3
3	5	5
4	4	6

- What if the kernel is separable?

– $O(n^2m) = O(2n^2m)$

1	2	1
---	---	---

 \ast

2	3	3
3	5	5
4	4	6

 $=$

	11	
	18	
	18	

1
2
1

 \ast

	11	
	18	
	18	

 $=$

	65	

Separability of Filters

- Is this separable? If yes, what is the separable version?

$$\frac{1}{K^2}$$

1	1	...	1
1	1	...	1
\vdots	\vdots	1	\vdots
1	1	...	1

$$\frac{1}{K}$$

1	1	...	1
---	---	-----	---

- What does this filter do?

Separability of Filters

- Is this separable? If yes, what is the separable version?

$$\frac{1}{256}$$

1	4	6	4	1
4	16	24	16	4
6	24	36	24	6
4	16	24	16	4
1	4	6	4	1

$$\frac{1}{16}$$

1	4	6	4	1
---	---	---	---	---

- What does this filter do?

Separability of Filters

- Is this separable? If yes, what is the separable version?

 $\frac{1}{8}$

-1	0	1
-2	0	2
-1	0	1

 $\frac{1}{2}$

-1	0	1
----	---	---

- What does this filter do?

Separability of Filters

- Is this separable? If yes, what is the separable version?

 $\frac{1}{4}$

1	-2	1
-2	4	-2
1	-2	1

 $\frac{1}{2}$

1	-2	1
---	----	---

- What does this filter do?

Image Sharpening

Intuition of Filtering



Original

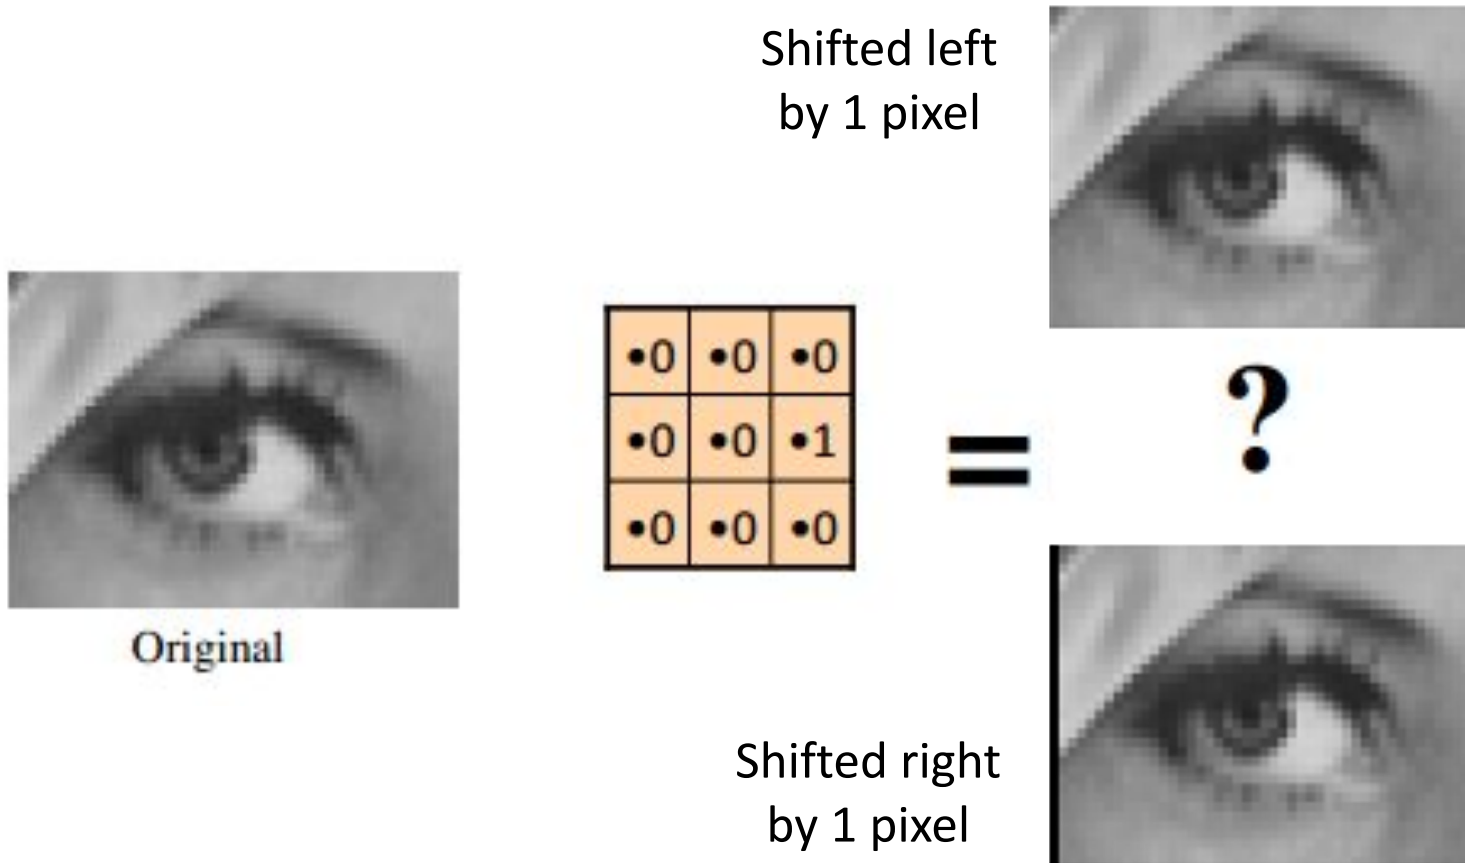
•0	•0	•0
•0	•1	•0
•0	•0	•0

=



Filtered
(no change)

Intuition of Filtering



Intuition of Filtering



Original

$$\frac{1}{9} \begin{bmatrix} \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \\ \bullet 1 & \bullet 1 & \bullet 1 \end{bmatrix} =$$



Blur (with a
box filter)

Sharpening

- **Sharpening by filtering:** Accentuates differences with local average



Original

$$* \frac{1}{9} \begin{bmatrix} -1 & -1 & -1 \\ -1 & 17 & -1 \\ -1 & -1 & -1 \end{bmatrix} =$$




$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

- This is also related to a process called **Unsharp masking**

Unsharp masking

- A process used many years in the publishing industry to sharpen images
- **Details = Original image – Smoothened image**
- **Sharpened image = Original image + Details**



Getting a sharpened image

- What does blurring take away?



- Adding the details back...



Sharpening Filter



Original

•0	•0	•0
•0	•2	•0
•0	•0	•0

-

$\frac{1}{9}$

•1	•1	•1
•1	•1	•1
•1	•1	•1

= ?

(Note that filter sums to 1)

“details of the image”

$$\begin{array}{|c|c|c|} \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \bullet 0 & \bullet 1 & \bullet 0 \\ \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \end{array} + \begin{array}{|c|c|c|} \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \bullet 0 & \bullet 1 & \bullet 0 \\ \hline \bullet 0 & \bullet 0 & \bullet 0 \\ \hline \end{array} - \frac{1}{9} \begin{array}{|c|c|c|} \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \bullet 1 & \bullet 1 & \bullet 1 \\ \hline \end{array}$$

The second 3x3 grid in the equation is highlighted with a red bracket and labeled “details of the image”.

Amount of sharpening

- **Details = Original image – Smoothened image**
- **Sharpened image = Original image + Details**
 - How can we control the amount of sharpening that is applied?

The strength of the details/smoothing filter!


Non-linear Filters

Non-linear filtering

- So far, we look at **linear filtering**

$$G[i, j] = \sum_{u=-k}^k \sum_{v=-k}^k H[u, v] F[i - u, j - v]$$

$$G = H \star F$$




Linear
combination of
multiplied terms

- What about **non-linear filtering**?
- Can the choice of filter $H[u, v]$ produce non-linear filtering?

Median Filter




- A type of **order-statistics** filter – “statistical estimator”
- No new pixel values are introduced




- Removes spikes: Good for impulse, salt & pepper noise

Median Filter

- Stray white pixels (very high values) or stray black pixels (very low values) can be dealt with



Median Filter



- Example:
 - Applying 7x7 median filter (middle pic) : **broccoli loses details**
 - Applying 7x7 multi-stage median filter : **less smoothing occurs, some details are maintained**

Median Filter

- **Multi-stage median filter:**
 - Median of a set of different median filters (obtained in different neighbourhoods)

$$y_{ij} = \text{med}(z_1, z_2, z_3, z_4)$$

$$z_1 = \text{med}(\{x_{uv} | x_{uv} \in N_{ij}^1\})$$

$$z_2 = \text{med}(\{x_{uv} | x_{uv} \in N_{ij}^2\})$$

$$z_3 = \text{med}(\{x_{uv} | x_{uv} \in N_{ij}^3\})$$

$$z_4 = \text{med}(\{x_{uv} | x_{uv} \in N_{ij}^4\})$$

- Preserves sharp corners

Median Filter


- Drawback: Tend to be better at rejecting **outliers** but not so good at handling noise without outliers (e.g. Gaussian white noise)
- Averaging filter / weighted average filter
 - too much blurriness, noise remains but smoothed

Alpha-trimmed mean filter

- Select the average of the values within a window which excludes a percent of the largest and smallest values in the neighbourhood
- Delete $\alpha/2$ lowest and $\alpha/2$ highest values in the neighbourhood S_{XY} , find average of remaining pixels:

$$\hat{f}(x, y) = \frac{1}{mn - \alpha} \sum_{(s,t) \in S_{XY}} g_r(s, t)$$

Alpha-trimmed mean filter



- How should α be chosen?

Summary

- **Linear filtering – Convolution**

- Smoothing by filtering
- Sharpening by filtering
- Unsharp masking

- **Non-linear filtering**

- Median filter
- Alpha-trimmed mean filter

- Next: Edge detection

Recommended Reading

- [Gonzalez & Woods] Chapter 3
- [Forsyth & Ponce] Chapter 8 & 10