

TDS3651

Visual Information Processing



Edges Lecture 4

Faculty of Computing and Informatics
Multimedia University

Lecture Outline

- What makes an edge?
- Gradient-based Edge Detectors
- Canny Edge Detector

Revision: Filtering

- What is the output of the following filters?

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

*

-1	0	1
-1	0	1
-1	0	1

=

Revision: Filtering

- What is the output of the following filters?

3 ⁻¹	0 ⁰	1 ¹	2	7	4
1 ⁻¹	5 ⁰	8 ¹	9	3	1
2 ⁻¹	7 ⁰	2 ¹	5	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

 \star

-1	0	1
-1	0	1
-1	0	1

 $=$

5			

Revision: Filtering

- What is the output of the following filters?

3	0 ⁻¹	1 ⁰	2 ¹	7	4
1	5 ⁻¹	8 ⁰	9 ¹	3	1
2	7 ⁻¹	2 ⁰	5 ¹	1	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

*

-1	0	1
-1	0	1
-1	0	1

=

5	4		

Revision: Filtering

- What is the output of the following filters?

3	0	1 ⁻¹	2 ⁰	7 ¹	4
1	5	8 ⁻¹	9 ⁰	3 ¹	1
2	7	2 ⁻¹	5 ⁰	1 ¹	3
0	1	3	1	7	8
4	2	1	6	2	8
2	4	5	2	3	9

*

-1	0	1
-1	0	1
-1	0	1

=

5	4	0	

Revision: Filtering

- What is the output of the following filters?

3	0	1	2	7	4
1	5	8	9	3	1
2	7	2	5	1	3
0	1	3	1 ⁻¹	7 ⁰	8 ¹
4	2	1	6 ⁻¹	2 ⁰	8 ¹
2	4	5	2 ⁻¹	3 ⁰	9 ¹

 $*$

-1	0	1
-1	0	1
-1	0	1

 $=$

5	4	0	-8
10	2	2	-3
0	2	4	7
3	2	3	16

Revision: Filtering

- What does the following filter do?

Filter 1

-1	0	1
-2	0	2
-1	0	1

Filter 2

1	-2	1
-2	4	-2
1	-2	1

- A. Smoothing
- B. Sharpening
- C. Extract details
- D. Denoising

Revision: Filtering

- What does the following filter do?

Filter 1

-1	0	1
-2	0	2
-1	0	1

Filter 2

1	-2	1
-2	4	-2
1	-2	1



- A. Smoothing
- B. Sharpening
- C. Extract details
- D. Denoising

Recall: Unsharp Masking



Original

$$* \frac{1}{9} \begin{bmatrix} -1 & -1 & -1 \\ -1 & 17 & -1 \\ -1 & -1 & -1 \end{bmatrix} =$$

Sum to 1
(normalized)

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Un-normalized filter

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} + \begin{bmatrix} 0 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 0 \end{bmatrix} - \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

Details of Image

Recall: Image filtering

- Compute a function of the local neighbourhood at each pixel in the image
 - Function specified by a “filter” or mask saying how to combine values from neighbours
- Uses of filtering
 - Enhance an image (denoising, sharpening, etc)
 - **Extract information (edges, textures, etc)**

What makes an edge?

Edge Detection

- **Goal:** Identify sudden changes (discontinuities) in an image
 - Most semantic and shape information from the image can be encoded in the edges
- **Ideal:** Artist's line drawing (but artist is also using object-level knowledge)



Source: D. Lowe

Edge Detection

- **Goal:** Identify sudden changes (discontinuities) in an image
 - ⇒ Map image from 2D array of pixels to a set of curves or line segments/contours
- **Main Idea:** Look for strong gradients, post-process


A close lookout for “edges”



A close lookout for “edges”



Surface normal discontinuity



Source: D. Holem

A close lookout for “edges”



Depth discontinuity



Source: D. Hoiem

A close lookout for “edges”



Surface color discontinuity




Source: D. Holem

A close lookout for “edges”



Potential edges detected!




Illumination discontinuity ?

Gradient/Derivative-based edge detectors

Derivatives and Edges

- An edge is a place of rapid change in the image intensity function



Derivatives with convolution

- For 2D function $f(x, y)$, the partial derivative:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

- For discrete data, we can approximate using finite differences:


$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x+1, y) - f(x, y)}{1}$$

- To implement it by **convolution**, what would be the **associated filter**?

Partial derivatives of an image




$$\frac{\partial f(x, y)}{\partial x}$$



-1	1
----	---

$$\frac{\partial f(x, y)}{\partial y}$$






-1	or	1
1		-1

Image gradient

- The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

- Gradient – Points in the direction of most rapid increase in intensity



- Direction:**

$$\theta = \tan^{-1} \left(\frac{\partial f / \partial y}{\partial f / \partial x} \right)$$

- Magnitude (“edge strength”)** $\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$


Partial derivatives of an image

- Taking both partial derivatives in the x direction and y direction, we get the full set of derivatives or “gradients” corresponding to both directions




Looking at a row profile

Intensity profile



Effects of noise



- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal



Where's the edge??

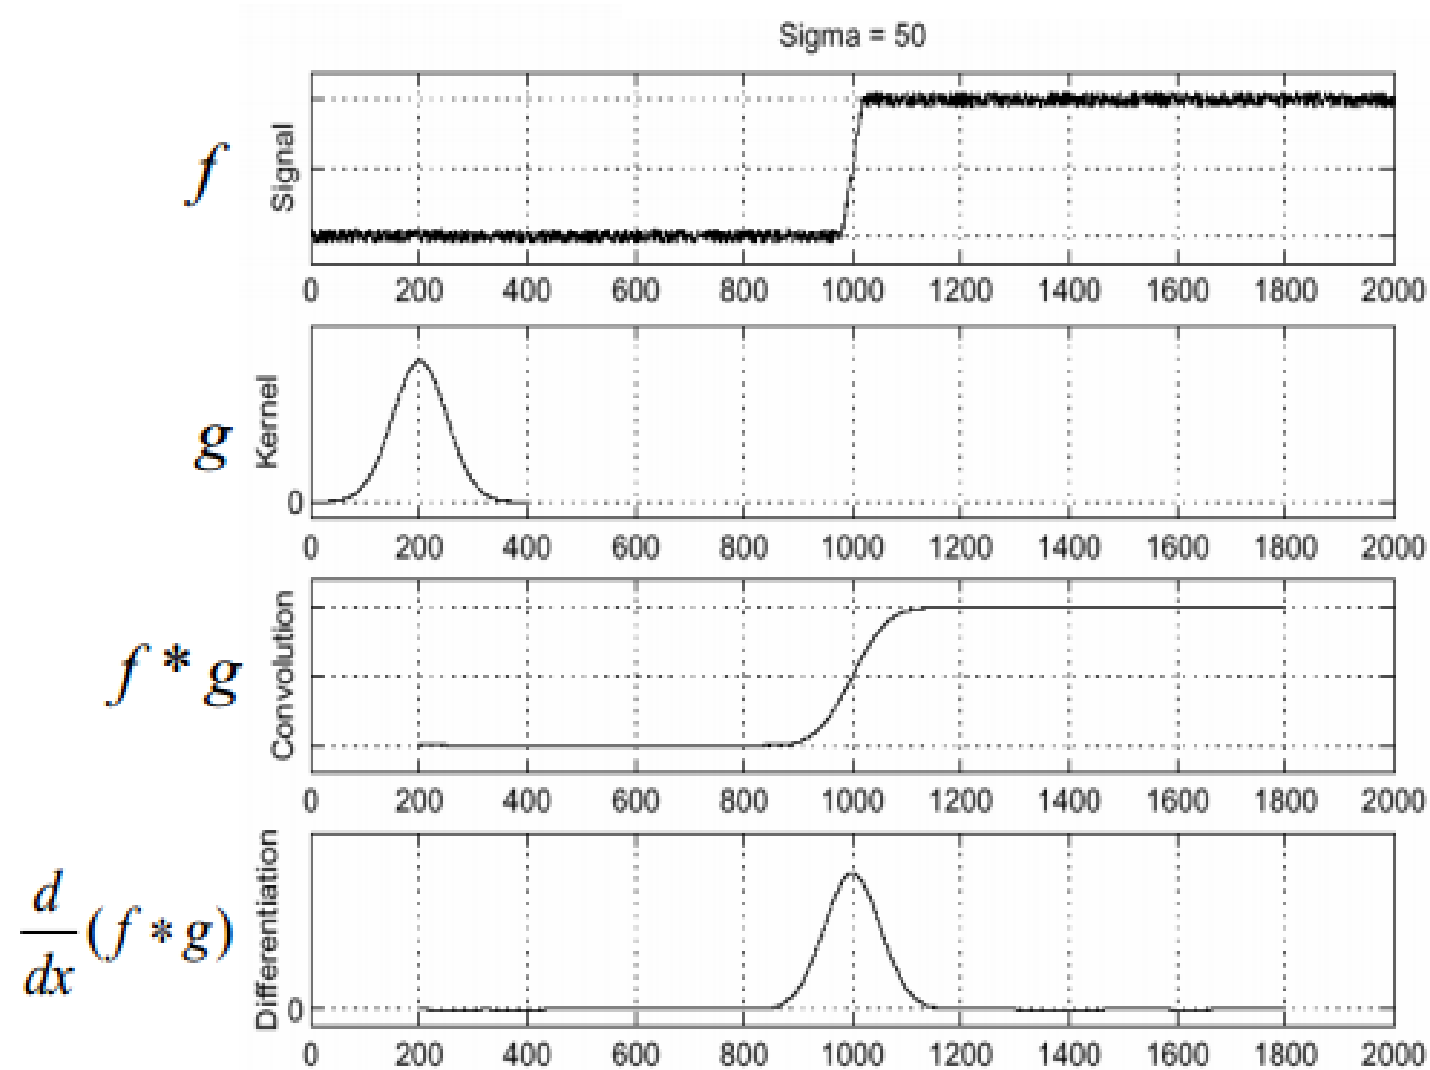
Effects of noise

- Gradient filters respond strongly to noise
 - Image noise results in pixels that look very different from their neighbours
 - What can be done?



- A. Smoothing to make pixels look more like their neighbours
- B. Thresholding to remove all the noise pixels
- C. Filtering to transform the pixels into a clearer format
- D. Nothing can be done

Solution:




Derivative theorem of convolution


- Differentiation is convolution, convolution is associative

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$


f



$\frac{d}{dx}g$




$f * \frac{d}{dx}g$




Saves one operation!

Derivative of Gaussian filter



$$* [1 \ -1] =$$




Derivative of Gaussian filter



$$(I \otimes g) \otimes h = I \otimes (g \otimes h)$$


$$\begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix} \otimes \begin{bmatrix} 1 & -1 \end{bmatrix}$$



Using 1st order gradient filter




**Recall, 1st
derivative
of Gaussian**




- Gradient magnitude is larger along a thick ridge
 - So, how to identify the actual edge points?
 - If we can, then how do we link the actual edge points to form curves?

Laplacian of Gaussian




- 2nd derivative of the Gaussian
- Where is the edge? **Zero-crossings of last graph**

2D Gradient-based Edge Filters




Gaussian

$$h_{\sigma}(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_{\sigma}(u, v)$$



Laplacian of Gaussian


$$\nabla^2 h_{\sigma}(u, v)$$

∇^2 Laplacian operator

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$


Laplacian of Gaussian (LoG) filter

- **LoG filter** – also known as **Marr-Hildreth** algorithm
 - Convolution of the image with the Laplacian of the Gaussian function
 - Zero crossings are detected to obtain edges
 - Its operator sometimes known as **Mexican hat operator**





DoG \Rightarrow LoG

- **Difference of Gaussians (DoG) filter** – a fast approximation of the LoG filter
 - Subtract one blurred version of original image (by Gaussian filter) from another, less blurred version of original
 - \Rightarrow Just use different σ values to create the two different blurred versions!




Smoothing with a Gaussian

- **Recall:** Parameter σ is the “scale” or “width” of the Gaussian kernel, controls amount of smoothing




...



Effect of σ on derivatives

- Edge structure differs depending on Gaussian kernel's scale parameter



- **Larger** values: thick, only important edges detected
- **Smaller** values: finer edges detected, but too much details
- So, what scale to choose?

Other **finite difference** filters

- “Finite difference filters” – approximates the gradient of image intensity function



Roberts

$$\begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$$



Prewitt

$$\begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$$

$$\begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$



Sobel

$$\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

$$\begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

Other finite difference filters

- “Finite difference filters” – approximates the gradient of image intensity function

0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

*

1	0	-1
1	0	-1
1	0	-1

=



Other finite difference filters

- “Finite difference filters” – approximates the gradient of image intensity function


0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10
0	0	0	10	10	10

*


1	0	-1
1	0	-1
1	0	-1

=

0	-30	-30	0
0	-30	-30	0
0	-30	-30	0
0	-30	-30	0



Other finite difference filters



Mask properties

- **Smoothing**


- Values positive
- Sum to 1 \Rightarrow constant regions same as input
- Amount of smoothing proportional to mask size

- **Edge Derivatives**

- Different signs (positive, negative) are used to get high response in regions of high contrast
 - **Sum to ?** \Rightarrow no response in constant regions
 - High absolute value at points of high contrast
- A. 3
B. 2
C. 1
D. 0


Designing an edge detector

- Criteria of an “optimal” edge detector
 - **Good detection**: Must minimize probability of false positives (spurious edges due to noise), and false negatives (missing real edges)



Designing an edge detector

- Criteria of an “optimal” edge detector
 - **Good detection:** Must minimize probability of false positives (spurious edges due to noise), and false negatives (missing real edges)
 - **Good localization:** Edges detected are as close as possible to true edges




Designing an edge detector

- Criteria of an “optimal” edge detector
 - **Good detection:** Must minimize probability of false positives (spurious edges due to noise), and false negatives (missing real edges)
 - **Good localization:** Edges detected are as close as possible to true edges
 - **Single response:** Detector must return one point only for each true edges




True
edge



Poor robustness
to noise



Poor
localization



Too many
responses

Designing an edge detector

- Primary edge detection steps:
 - **Smoothing:** Suppress noise
 - **Edge enhancement:** Filter for contrast
 - **Edge localization:** Determine which local maxima from filter output are actually edges
 - **Thresholding**

Thresholding

- Choose a threshold value t
- Set any pixels less than t to **zero** (off)
- Set any pixels greater or equal to t to **one** (on)

Original image



Thresholding: Gradient magnitude image



Which threshold to use?



Lower threshold value



Higher threshold value

What can we conclude from using different threshold values?

- A. A universal value can be used for the best thresholding.
- B. Lower threshold will detect all edges we need.
- C. Higher threshold may miss out on some edges.
- D. A random value can be used to perform thresholding.

Canny edge detector

Canny edge detector

- Most widely used edge detector – Stable, consistent
- Steps:

1. **Filter image with Gaussian filter**, find magnitude and orientation of gradient

2. **Non-maximum suppression**

- Thin wide “ridges” down to single pixel width (get rid of spurious responses)

3. **Hysteresis thresholding and linking**

Normally we just threshold after filtering

- Define 2 thresholds low and high (to mark strong and weak edges)
- Start edge curves from the true edges and continue tracking based on weak edges that are connected to the true edges


Canny edge detector

- Most widely used edge detector – Stable, consistent
- Steps:
 - 1. Filter image with Gaussian filter**, find magnitude and orientation of gradient
 - 2. Non-maximum suppression**
 - Thin wide “ridges” down to single pixel width (get rid of spurious responses)
 - 3. Hysteresis thresholding and linking**
 - Define 2 thresholds: low and high (to mark strong and weak edges)
 - Start edge curves from the true edges and continue tracking based on weak edges that are connected to the true edges

Step 1




- Lena Example:
 - Step 1 can be done in one pass, by filtering with the derivative of Gaussian filter (both x and y directions)



2 Problems when thresholding

- Problem #1:
 - Edges from many gradient-based edge detectors are too thick



How to turn these thick regions into single pixel edges?

Canny edge detector

- Most widely used edge detector – Stable, consistent
- Steps:

1. **Filter image with Gaussian filter**, find magnitude and orientation of gradient

2. **Non-maximum suppression**

- Thin wide “ridges” down to single pixel width (get rid of spurious responses)


3. **Hysteresis thresholding and linking**

- Define 2 thresholds: low and high (to mark strong and weak edges)
- Start edge curves from the true edges and continue tracking based on weak edges that are connected to the true edges

Step 2



- Get orientation at each pixel

$$\theta = \tan^{-1} \left(\frac{\partial f / \partial y}{\partial f / \partial x} \right)$$



- **Non-maximum suppression**

- Check if pixel is local maximum along gradient direction e.g. maximum at q if value larger than p and r (interpolated pixels). So, suppress p and r .



2 Problems when thresholding

- Problem #2:
 - Pixels along weak edges will not survive if standard thresholding is applied (using a threshold determined globally)



How can we keep the pixels
along

- A. The
- B. Mul
- C. The
- D. All t



keep them.
n them.
weak.
ense.

[Result](#)

Canny edge detector

- Most widely used edge detector – Stable, consistent
- Steps:

1. **Filter image with Gaussian filter**, find magnitude and orientation of gradient

2. **Non-maximum suppression**


- Thin wide “ridges” down to single pixel width (get rid of spurious responses)

3. **Hysteresis thresholding and linking**

- Define 2 thresholds: low and high (to mark strong and weak edges)
- Start edge curves from the true edges and continue tracking based on weak edges that are connected to the true edges


Step 3

- Double threshold
 - **High threshold** – (Gradient higher than this threshold) \Rightarrow **strong** edges
 - **Low threshold** – (Gradient higher than low threshold but lower than high threshold) \Rightarrow **weak** edges
- The plan: Ensure continuity in edges of different strengths




Step 3


- To “link” edges, perform connected components, starting from strong edge pixels
 - Search 8-neighbourhood for weak edge pixels
 - Usually a weak edge pixel caused by true edges will be connected to a strong edge pixel while noise responses are unconnected.



Strong edge pixel



Weak edge pixel




Linked edges



Final Result



Effect of σ (Gaussian kernel width)



- The choice of σ depends on desired behaviour
 - Large σ detect large scale edges (less edges typically)
 - Small σ detects fine features (more edges typically)


So...how do we know if the edges found are really the **TRUE** edges?

Gradient edges vs. Human segmentation



image



human segmentation



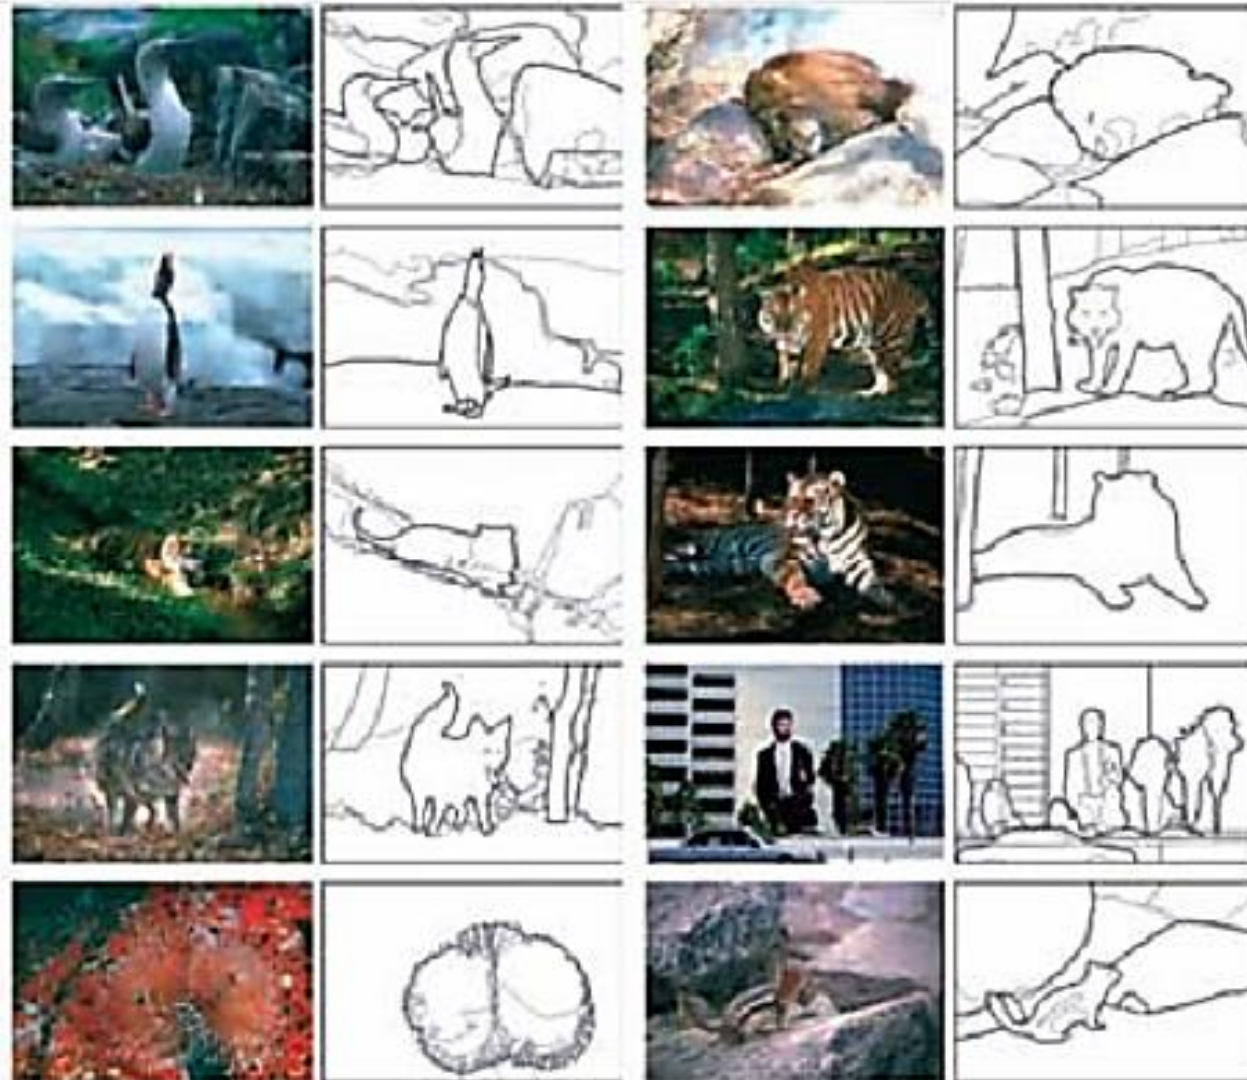
gradient magnitude



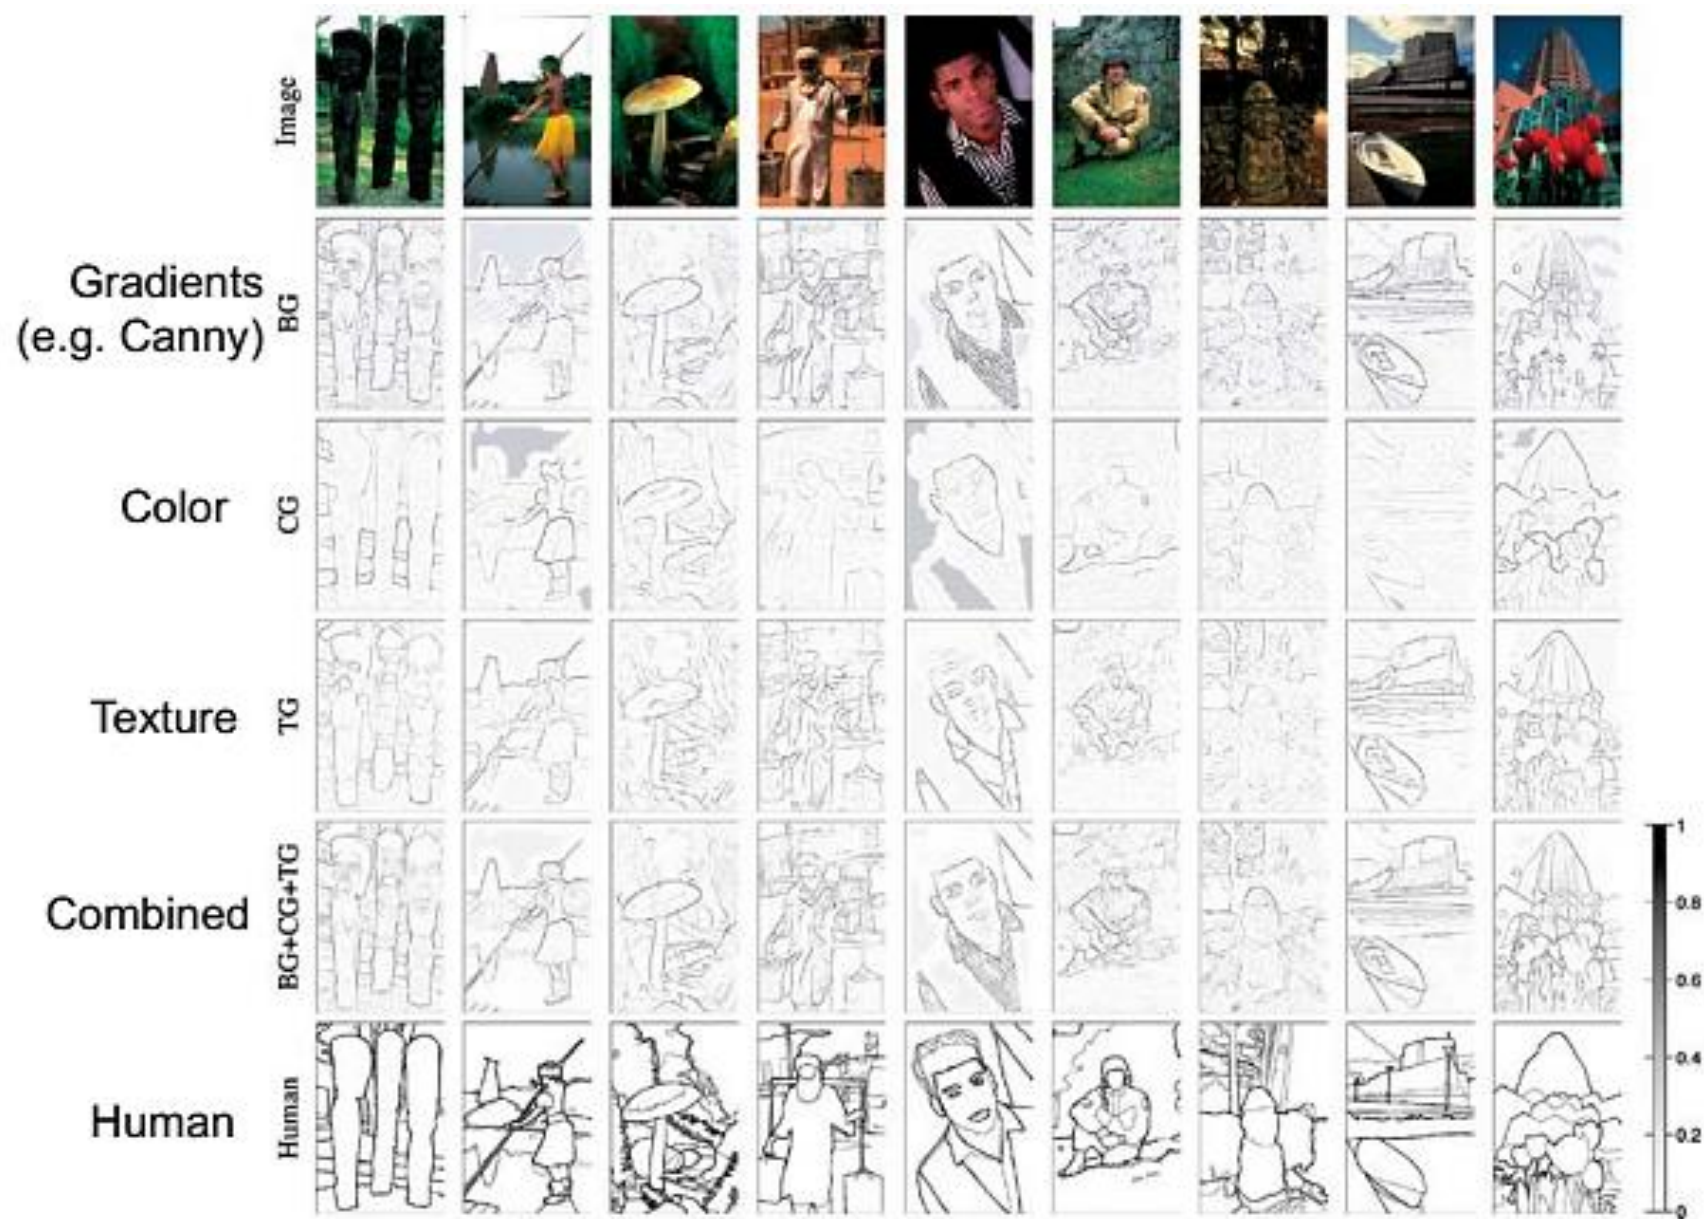
Berkeley segmentation database

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>


Human-perceived edges




Features used vs. Human-marked



Contour Detection (CVPR 2008)



Contour Detection (ICCV 2013)



Best performing:

P. Dollar and C. Zitnick, “**Structured Forests for Fast Edge Detection**” (ICCV 2013)

Code:

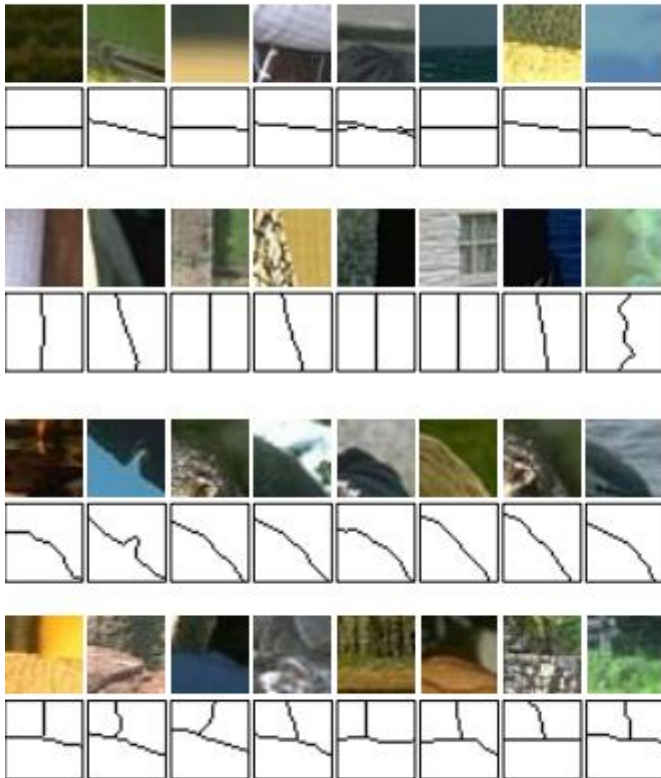
Original (Matlab/C++): <https://www.microsoft.com/en-us/download/details.aspx?id=52370>

Python: <https://github.com/ArtanisCV/StructuredForests>

OpenCV (C++): https://docs.opencv.org/3.3.1/d0/da5/tutorial_ximgproc_prediction.html

Data-driven edge detection


Training data



Input images



Ground truth



Output



Machine learning: How can we “train” an edge detector to learn what is an edge?

Deep Learning-based Edge detection

State-of-the-arts deep-learning based edge detection:

- **Holistically-Nested Edge Detection**

<https://ieeexplore.ieee.org/abstract/document/7410521/>

- **Learning Relaxed Deep Supervision for Better Edge Detection**

<https://ieeexplore.ieee.org/document/7780401>

- **DeepContour: A deep convolutional feature learned by positive-sharing loss for contour detection**

<https://ieeexplore.ieee.org/document/7299024/>

Summary

- **Edges – what makes an edge?**
- **Gradient-based edge detectors**
 - Using derivatives or gradients to find edges
 - LoG and DoG
- **Canny edge detector**
- **Knowing what are true edges**

Recommended Reading

- [Gonzalez & Woods] Chapter 10
- [Forsyth & Ponce] Chapter 8